Jing-Jing Li

jl3676@berkeley.edu
http://jl3676.github.io
Google Scholar

Education

University of California, Berkeley — Berkeley, CA
Ph.D. in Neuroscience with concentrations in Computation and Cognition, GPA: 3.97/4.00 — 2021–2026

Cornell University — Ithaca, NY
B.A. in Computer Science and Mathematics, Minor in Cognitive Science, GPA: 4.01/4.30 — 2017–2020

Work Experiences

Amazon Web Services — Seattle, WA
Applied Scientist Intern — May 2025 - August 2025

- Will conduct AI safety research for AWS Bedrock.

Allen Institute for Artificial Intelligence — Seattle, WA
Research Intern — May 2024 - August 2024

- Developed a framework to improve the interpretability, transparency, and steerability of AI safety moderation.
- Performed prompt engineering, taxonomy development, crowdsourcing, model distillation, supervised fine-tuning, and evaluation on LLMs.

Skills

- **Large Language Models:** Prompt Engineering, Supervised Fine-Tuning, Crowdsourcing
- **Machine Learning:** Pytorch, TensorFlow, scikit-learn, CUDA, Hugging Face
- **Programming:** Python, Java, C, C++, Bash, Shell, HTML, CSS, JavaScript, GitHub
- **Data Science:** Numpy, SciPy, pandas, Matplotlib, R, MATLAB, SQL
- **Other:** LaTeX, Adobe Illustrator, Adobe Photoshop, Linux, Microsoft Excel

Relevant Courses

- **Machine Learning:** Deep Unsupervised Learning, LLMs and Cognition, Deep Reinforcement Learning, Computer Vision, Large-Scale Machine Learning, Intro to Machine Learning, Computational Genetics
- **Software Engineering:** Data Structures (Honors), Algorithms, Operating Systems, Database Systems
- **Mathematics and Statistics:** Numerical Analysis, Biological Statistics, Probability Theory, Abstract Algebra, Linear Algebra (Honors), Multi-variable Calculus

Grants and Fellowships

- UC Berkeley ICBS Grant ($5,000; *Co-recipient with Eve Fleisig*) — 2024-2025
- Society for Neuroscience Trainee Professional Development Award — 2024
- CogSci Conference Travel Grant — 2023
- Milton I. and Florence Mack Neurology Research Fund — 2021–2022
- Summer Undergraduate Research Fellowship, Caltech — 2018

PUBLICATIONS

[1] **J.-J. Li**, V. Pyatkin, M. Kleiman-Weiner, L. Jiang, N. Dziri, A. G. E. Collins, J. S. Borg, M. Sap, Y. Choi, and S. Levine, "SafetyAnalyst: Interpretable, transparent, and steerable safety moderation for AI behavior", in *International conference on machine learning*, PMLR, 2025.

[2] **J.-J. Li**, C. Chen, and A. G. Collins, "Humans integrate heuristics and bayesian inference to efficiently explore under uncertainty", in *Proceedings of the Annual Meeting of the Cognitive Science Society*, 2025.

[3] **J.-J. Li** and A. G. Collins, "An algorithmic account for how humans efficiently learn, transfer, and compose hierarchically structured decision policies", *Cognition*, vol. 254, p. 105 967, 2025.

[4] J. Chase, **J.-J. Li**, W. C. Lin, L.-H. Tai, A. G. Collins, and L. Wilbrecht, "Genetic changes linked to two different syndromic forms of autism enhance reinforcement learning in adolescent male but not female mice", *bioRxiv*, pp. 2025–01, 2025.

[5] **J.-J. Li**, C. Shi, L. Li, and A. G. Collins, "Dynamic noise estimation: A generalized method for modeling noise fluctuations in decision-making", *Journal of Mathematical Psychology*, vol. 119, p. 102 842, 2024.

[6] T.-F. Pan, **J.-J. Li**, B. Thompson, and A. Collins, *Latent variable sequence identification for cognitive models with neural bayes estimation*, 2024. arXiv: 2406.14742 `[cs.LG]`.

[7] D. S. Jin, O. Agdali, T. Yadav, S. I. Kronemer, S. Kunkler, S. Majumder, M. Khurana, M. C. McCusker, I. Fu, A. Khalaf, K. L. Christison-Lagay, S. L. Aerts, Q. Xin, **J.-J. Li**, S. H. McGill, M. J. Crowley, and H. Blumenfeld, "Neural mechanisms of awareness of action", *bioRxiv*, 2024.

[8] **J.-J. Li**, C. Shi, L. Li, and A. G. Collins, "A generalized method for dynamic noise inference in modeling sequential decision-making", in *Proceedings of the Annual Meeting of the Cognitive Science Society*, 2023.

[9] C. McCafferty, B. F. Gruenbaum, R. Tung, **J.-J. Li**, X. Zheng, P. Salvino, P. Vincent, Z. Kratochvil, J. H. Ryu, A. Khalaf, K. Swift, R. Akbari, W. Islam, P. Antwi, E. A. Johnson, P. Vitkovskiy, J. Sampognaro, I. G. Freedman, A. Kundishora, A. Depaulis, F. David, V. Crunelli, B. G. Sanganahalli, P. Herman, F. Hyder, and H. Blumenfeld, "Decreased but diverse activity of cortical and thalamic neurons in consciousness-impairing rodent absence seizures", *Nature Communications*, vol. 14, no. 1, pp. 1–19, 2023.

[10] **J.-J. Li**, L. Xia, F. Dong, and A. G. Collins, "Credit assignment in hierarchical option transfer", in *Proceedings of the Annual Meeting of the Cognitive Science Society*, 2022.

[11] J. Ding, **J.-J. Li**, and M. Xu, "Classification of murmurs in pcg using combined frequency domain and physician inspired features", in *2022 Computing in Cardiology (CinC)*, IEEE, vol. 498, 2022, pp. 1–4.

PRESENTATIONS

Invited talks

Dynamic noise modeling in decision-making	Uniklinikum Würzburg
Cognitive and Computational Neuroscience in Development Psychiatry Research Group	June 2024

Conference talks

Humans integrate heuristics and Bayesian inference to efficiently explore	San Francisco, CA
CogSci Conference	July 2025
Dynamic noise modeling in decision-making	Tahoe, CA
Berkeley Neuroscience Conference	October 2023
A generalized method for dynamic noise inference	Sydney, Australia
CogSci Conference	July 2023

Credit assignment in the transfer of hierarchical options	Toronto, Canada
CogSci Conference	July 2022

Conference posters

Interpretable, transparent, and steerable LLM safety moderation	Vancouver, Canada
ICML	July 2025
🔦 **Humans Integrate heuristics and Bayesian inference to efficiently explore**	Dublin, Ireland
RLDM Conference (spotlight)	June 2025
Interpretable, transparent, and steerable LLM safety moderation	Vancouver, Canada
NeurIPS SoLaR Workshop	December 2024
Modeling how humans learn, transfer, and compose hierarchical policies	Chicago, IL
Society for Neuroscience Conference	October 2024
Modeling the emergence of instrumental learning in an odor-based 2AFC task	Boston, MA
Cognitive Computational Neuroscience Conference	August 2024
Modeling how humans learn, transfer, and compose hierarchical policies	Boston, MA
Cognitive Computational Neuroscience Conference	August 2024
Credit assignment in the learning and transfer of hierarchical options	San Francisco, CA
Cognitive Neuroscience Society Conference	April 2022